

Chapter 1

Fun with Floats

Floats are inexact by nature and this can confuse programmers. In this chapter we present an introduction to this problem. The basic message is that Floats are what they are inexact but fast numbers.

1.1 Never test equality on floats

The first basic principle is to never compare float equality. Let's take a simple case: the addition of two floats may not be equal to the float representing their sum. For example $0.1 + 0.2$ is not equal to 0.3 .

```
(0.1 + 0.2) = 0.3  
returns false
```

Hey, this is unexpected, you did not learn that in school, did you? This behavior is surprising indeed, but it's normal since floats are inexact numbers. What is important to understand is that the way floats are printed is also impacting our understanding. Some approaches prints a simpler representation of reality than others. In early versions of Pharo printing $0.1 + 0.2$ was printing 0.3 , now printing it returns 0.30000000000000004 . This change was guided by the idea that it is better not to lie to the user. Showing the inexactness of float is better than hiding because one day or another we can be deeply bitten by them.

```
(0.2 + 0.1 ) printString  
returns '0.30000000000000004'
```

```
0.3 printString  
returns '0.3'
```

We can see that we are in presence of two different numbers by looking at the hexadecimal values.

```
(0.1+0.2) hex
returns '3FD3333333333334'
0.3 hex
returns '3FD3333333333333'
```

The method `storeString` also conveys that we are in presence of two different numbers.

```
(0.1+0.2) storeString
returns '0.30000000000000004'
0.3 storeString
returns '0.3'
```

About `closeTo`: One way to know if two floats are probably close enough to look like the same number is to use the message `closeTo`:

```
(0.1 + 0.2) closeTo: 0.3
returns true

0.3 closeTo: (0.1 + 0.2)
returns true
```

About Scaled Decimals. Scaled Decimals are exact numbers so they exhibit the behavior you expected.

```
0.1s2 + 0.2s2 = 0.3s2
returns true
```

Analyzing more deeply above example

To understand what operation is involved in above addition, we must know how Float are represented internally in the computer: Pharo's Float format is a wide spread standard found on most computers - IEEE 754-1985 double precision on 64 bits (See http://en.wikipedia.org/wiki/IEEE_754-1985 for more details). In this format, a Float is represented in base 2 by this formula:

$$sign \cdot mantissa \cdot 2^{exponent}$$

- The sign is represented on 1 bit.
- The exponent is represented on 11 bits.

- The mantissa is a fractional number in base two, with a leading 1 before decimal point, and with 52 binary digits after fraction point.

For example, a serie of 52 bits:

011001000

means the mantissa is:

1.011001000

which also means

$$1 + \frac{0}{2} + \frac{1}{2^2} + \frac{1}{2^3} + \frac{0}{2^4} + \frac{0}{2^5} + \frac{1}{2^6} + \cdots + \frac{0}{2^{52}}$$

The mantissa is thus between 1 (included) and 2 (excluded) for normal numbers.

Let us construct such a mantissa:

```
(#(0 2 3 6) detectSum: [:i | (2 raisedTo: i) reciprocal]) asFloat.  
returns 1.390625
```

Now let us multiply by 2^3 to get a non null exponent:

```
(#(0 2 3 6) detectSum: [:i | (2 raisedTo: i) reciprocal]) asFloat timesTwoPower: 3.  
returns 11.125
```

In Pharo, you can retrieve these informations:

11.125 sign.
returns 1

11.125 significand.
returns 1.390625

11.125 exponent.
returns 3

Generally you won't handle the mantissa, but rather the mantissa shifted 52 bits to the left in order to operate on exact integers. The result includes the leading 1 and should thus be 53 bits long for a normal number (that's the float precision):

```
11.125 significandAsInteger printStringBase: 2.  
returns '1011001000000000000000000000000000000000000000000000000000000000'
```

```
11.125 significandAsInteger highBit.  
    returns 53
```

Float precision.
returns 53

You can also retrieve the exact fraction corresponding to the internal representation of the Float:

```
11.125 asTrueFraction.  
returns (89/8)
```

```
(#(0 2 3 6) detectSum: [:i | (2 raisedTo: i) reciprocal]) * (2 raisedTo: 3).  
returns (89/8)
```

Until there we retrieved the exact input we've injected into the Float. Are Float operations exact after all? Hem, no, we only played with fractions having a power of 2 as denominator and a few bits in numerator. If one of these conditions is not met, we won't find any exact Float representation of our numbers. In particular, it is not possible to represent $1/5$ with a finite number of binary digits. Consequently, a decimal fraction like 0.1 cannot be represented exactly with above representation.

```
(1/5) asFloat = (1/5).  
returns false
```

```
(1/10) = 0.1  
returns false
```

Let us see in detail how we could get the fractional bits of $2r1 / 2r101$: we must pose the division:

1		101
10		0.00110011
100		
1000		
-101		
11		
110		
-101		
1		
10		
100		
1000		
-101		
11		
110		
-101		
1		

What we see is that we get a cycle: every 4 Euclidean divisions, we get a quotient $2r0011$ and a remainder 1. That means that we need an infinite serie

1.3 Another example: rounding is inexact

While float equality is known to be evil, you have to pay attention to other aspects of floats. Let us illustrate that point with the following example.

```
2.8 truncateTo: 0.01  
returns 2.8000000000000003
```

```
2.8 roundTo: 0.01  
returns 2.8000000000000003
```

It is surprising but not false that `2.8 truncateTo: 0.01` does not return `2.8` but `2.8000000000000003`. This is because `truncateTo:` and `roundTo:` perform several operations on floats: inexact operations on inexact numbers can lead to cumulative rounding errors as you saw above, and that's just what happens again.

Even if you perform the operations exactly and then round to nearest Float, the result is inexact because of the initial inexact representation of `2.8` and `0.01`.

```
(2.8 asTrueFraction roundTo: 0.01 asTrueFraction) asFloat  
returns 2.8000000000000003
```

Using `0.01s2` rather than `0.01` let this example appear to work:

```
2.80 truncateTo: 0.01s2  
returns 2.80s2
```

```
2.80 roundTo: 0.01s2  
returns 2.80s2
```

But it's just a case of luck, the fact that `2.8` is inexact is enough to cause other surprises as illustrated below:

```
2.8 truncateTo: 0.001s3.  
returns 2.799s3
```

```
2.8 < 2.800s3.  
returns true
```

Truncating in the Float world is absolutely unsafe. Though, using a `ScaledDecimal` for rounding is unlikely to cause such discrepancy, except when playing with last digits.

1.4 Fun with Inexact representations

To add a nail to the coffin, let's play a bit more with inexact representations. Let us try to see the difference between different numbers:

```
{
  ((2.8 asTrueFraction roundTo: 0.01 asTrueFraction) - (2.8 predecessor)) abs -> -1.
  ((2.8 asTrueFraction roundTo: 0.01 asTrueFraction) - (2.8)) abs -> 0.
  ((2.8 asTrueFraction roundTo: 0.01 asTrueFraction) - (2.8 successor)) abs -> 1.
} detectMin: [:e | e key ]
```

returns the pair
0.0->1

you get 0.0->1, which means that: $(2.8 \text{ asTrueFraction roundTo: } 0.01 \text{ asTrueFraction}) \text{ asFloat} = (2.8 \text{ successor})$

But remember that

```
(2.8 asTrueFraction roundTo: 0.01 asTrueFraction) ~= (2.8 successor)
```

It must be interpreted as the nearest Float to $(2.8 \text{ asTrueFraction roundTo: } 0.01 \text{ asTrueFraction})$ is (2.8 successor) .

If you want to know how far it is, then get an idea with:

```
((2.8 asTrueFraction roundTo: 0.01 asTrueFraction) - (2.8 successor asTrueFraction))
  asFloat
returns -2.0816681711721685e-16
```

1.5 Conclusion

Floats are inexact numbers. Pay attention when you manipulate them. There are much more things to know about floats, and if you are advanced enough, it would be a good idea to check this link from the wikipedia page "What Every Computer Scientist Should Know About Floating-Point Arithmetic".